

```

• begin
•   using Pkg
•   Pkg.activate(joinpath(Pkg.devdir(), "MLCourse"))
•   using CSV, DataFrames, Distributions, Plots, MLJ, MLJLinearModels, Random,
•       Statistics, OpenML
• end

```

# Linear Methods

We load the precipitation data from a csv file on the harddisk to a DataFrame. Our goal is to predict whether there is some precipitation (rain, snow etc.) on the next day in Pully, getting measurements from different weather stations in Switzerland.

```

• precipitation = CSV.read(joinpath(@__DIR__, "..", "data", "project",
  "trainingdata.csv"), DataFrame);

```

First we have to prepare our data set by dropping the missing values and split the datas into a train and a test set.

p =

	ABO_radiation_1	ABO_delta_pressure_1	ABO_air_temp_1	ABO_sunshine_1	ABO_win
1	-0.166667	-1.2	-5.68333	0.0	2.08333
2	0.333333	0.2	5.16667	0.0	1.43333
3	16.5	-0.3	6.16667	33.0	0.983333
4	5.33333	-0.6	9.01667	0.0	7.6
5	9.83333	0.3	7.38333	0.0	7.98333
6	25.5	-1.11022e-16	4.36667	48.0	0.483333
7	0.0	-1.0	-1.15	0.0	1.85
8	12.3333	0.8	13.45	20.0	0.516667
9	0.166667	-1.3	4.3	0.0	0.066666
10	0.333333	0.3	2.9	0.0	9.35
⋮ more					
1699	17.3333	0.1	14.9333	28.0	4.48333

```

• p = dropmissing(precipitation)

```

data\_split (generic function with 1 method)

```
• function data_split(data;
•                               shuffle = false,
•                               idx_train = 1:1275,
•                               idx_test = 1276:1699)
•   idxs = if shuffle
•           randperm(size(data, 1))
•         else
•           1:size(data, 1)
•         end
•   (train = data[idxs[idx_train], :],
•     test = data[idxs[idx_test], :])
• end
```

- `p1 = coerce!(p, :precipitation_nextday => Binary);` # with this we tell the computer to interpret the data in column precipitation\_nextday as multi-class data.

data1 =

```
►(train =
```

	ABO_radiation_1	ABO_delta_pressure_1	ABO_air_temp_1	ABO_sunshine_1
1	-0.166667	-1.2	-5.68333	0.0
2	0.333333	0.2	5.16667	0.0
3	16.5	-0.3	6.16667	33.0
4	5.33333	-0.6	9.01667	0.0
5	9.83333	0.3	7.38333	0.0
6	25.5	-1.11022e-16	4.36667	48.0
7	0.0	-1.0	-1.15	0.0
8	12.3333	0.8	13.45	20.0
9	0.166667	-1.3	4.3	0.0
10	0.333333	0.3	2.9	0.0
...	more			
1275	0.333333	-0.2	8.78333	0.0

```
• data1 = data_split(p1)
```

## Multiple Logistic Regression

Now we define a supervised learning machine.

```
• mach = machine(LogisticClassifier(penalty = :none),
•             select(data1.train, Not(:precipitation_nextday)),
•             data1.train.precipitation_nextday) |> fit!;
```

	Ground Truth	
Predicted	false	true
false	710	0
true	0	565

- `confusion_matrix(predict_mode(mach, select(data1.train, Not(:precipitation_nextday))), data1.train.precipitation_nextday)`

With our simple features, logistic regression can classify the training data almost always correctly. Let us see how well this works for test data.

	Ground Truth	
Predicted	false	true
false	180	40
true	69	135

- `confusion_matrix(predict_mode(mach, select(data1.test, Not(:precipitation_nextday))), data1.test.precipitation_nextday)`

Let us evaluate the fit in terms of commonly used losses for binary classification.

- `function losses(machine, input, response)`
- `(loglikelihood = -sum(log_loss(predict(machine, input), response)),`
- `misclassification_rate = mean(predict_mode(machine, input) != response),`
- `accuracy = accuracy(predict_mode(machine, input), response),`
- `auc = MLJ.auc(predict(machine, input), response)`
- `)`
- `end;`

► (loglikelihood = -9.35433, misclassification\_rate = 0.0, accuracy = 1.0, auc = 1.0)

- `losses(mach, select(data1.train, Not(:precipitation_nextday)), data1.train.precipitation_nextday)`

► (loglikelihood = -3464.55, misclassification\_rate = 0.257075, accuracy = 0.742925, auc

- `losses(mach, select(data1.test, Not(:precipitation_nextday)), data1.test.precipitation_nextday)`

## Multiple Logistic Ridge Regression

- `mach1 = machine(LogisticClassifier(penalty = :l2, lambda = 2e-2),`
- `select(data1.train, Not(:precipitation_nextday)),`
- `data1.train.precipitation_nextday) |> fit!;`

	Ground Truth	
Predicted	false	true
false	708	3
true	2	562

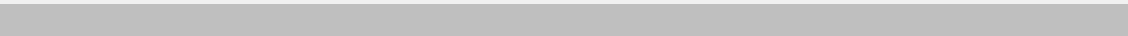
- `confusion_matrix(predict_mode(mach1, select(data1.train, Not(:precipitation_nextday))),`
- `data1.train.precipitation_nextday)`

	Ground Truth	
Predicted	false	true
false	182	44
true	67	131

- `confusion_matrix(predict_mode(mach1, select(data1.test,`
- `Not(:precipitation_nextday))),`
- `data1.test.precipitation_nextday)`

We see that the test misclassification rate with regularization is lower than in our original fit without regularization. The misclassification rate on the training set is higher. This indicates that unregularized logistic regression is too flexible for our data set.

► (loglikelihood = -33.0821, misclassification\_rate = 0.00392157, accuracy = 0.996078, au

◀  ▶

- `losses(mach1, select(data1.train, Not(:precipitation_nextday))),`
- `data1.train.precipitation_nextday)`

► (loglikelihood = -2439.7, misclassification\_rate = 0.261792, accuracy = 0.738208, auc =

◀  ▶

- `losses(mach1, select(data1.test, Not(:precipitation_nextday))),`
- `data1.test.precipitation_nextday)`