

```

• begin
•   using Pkg
•   Pkg.activate(joinpath(Pkg.devdir(), "MLCourse"))
•   using CSV, DataFrames, Distributions, Plots, MLJ, MLJLinearModels, Random,
•       Statistics, OpenML, MLJDecisionTreeInterface, MLJFlux, Flux
• end

```

## Non-Linear Methods

We load the precipitation data from a csv file on the harddisk to a DataFrame. Our goal is to predict whether there is some precipitation (rain, snow etc.) on the next day in Pully, getting measurements from different weather stations in Switzerland.

```

• precipitation = CSV.read(joinpath(@__DIR__, "..", "data", "project",
  "trainingdata.csv"), DataFrame);

```

```

• p = dropmissing(precipitation);

```

data\_split (generic function with 1 method)

```

• p1 = coerce!(p, :precipitation_nextday => Binary);

```

```

• data1 = data_split(p1);

```

## Tree-Based Methods

```

• mach = machine(RandomForestClassifier(n_trees = 500),
•       select(data1.train, Not(:precipitation_nextday)),
•       data1.train.precipitation_nextday);

```

```

• fit!(mach, verbosity = 2);

```

0.8136792452830188

```

• mean(predict_mode(mach, select(data1.test, Not(:precipitation_nextday))) .==
  data1.test.precipitation_nextday)

```

The test accuracy of a random forest with 500 trees is approximately 81%, better than with the linear methods (74% and ...).

Let's prepare these results for a submission data set. First we have to load the test set, and apply our machine on it. Then we construct our submission data and download it.

```

• md" Let's prepare these results for a submission data set. First we have to load the
  test set, and apply our machine on it. Then we construct our submission data and
  download it."

```

```
• precipitation_test = CSV.read(joinpath(@__DIR__, "..", "data", "project",  
"testdata.csv"), DataFrame);
```

```
pred =
```

```
► MLJBase.UnivariateFiniteVector{ScientificTypesBase.Multiclass{2}, Bool, UInt32, Float64
```

```
◀
```

```
• pred = predict(mach, precipitation_test)
```

```
true_pred =
```

```
► [0.96, 0.716, 0.794, 0.268, 0.776, 0.364, 0.118, 0.94, 0.212, 0.172, 0.356, 0.016, 0.018,
```

```
◀
```

```
• true_pred = pdf.(pred, true)
```

```
• submission = DataFrame(id = 1:1200, precipitation_nextday = true_pred);
```

```
"../data/project/submission_tree.csv"
```

```
• CSV.write("../data/project/submission_tree.csv", submission)
```

## Neural Networks

```
• mach1 = machine(NeuralNetworkClassifier(builder = MLJFlux.Short(n_hidden = 128,  
• dropout = 0.1,  
• σ = relu),  
• batch_size = 32,  
• epochs = 30),  
• select(data1.train, Not(:precipitation_nextday)),  
• data1.train.precipitation_nextday);
```

```
• fit!(mach1, verbosity = 2);
```

```
0.7735849056603774
```

```
• mean(predict_mode(mach1, select(data1.test, Not(:precipitation_nextday))) .==  
data1.test.precipitation_nextday)
```

```
pred1 =
```

```
► MLJBase.UnivariateFiniteVector{ScientificTypesBase.Multiclass{2}, Bool, UInt32, Float64
```

```
◀
```

```
• pred1 = predict(mach1, precipitation_test)
```

```
true_pred1 =
```

```
► [1.0, 1.0, 1.0, 0.0, 1.0, 1.0, 0.0, 1.0, 0.0, 0.0, 1.0, 0.0, 0.0, 0.0, 1.0, 1.0, 0.0, 1.0,
```

```
◀
```

```
• true_pred1 = pdf.(pred1, true)
```

```
• submission1 = DataFrame(id = 1:1200, precipitation_nextday = true_pred1);
```

```
"../data/project/submission_neural.csv"
```

```
• CSV.write("../data/project/submission_neural.csv", submission1)
```

